



School of Computing

Tutorial 7: BST and bBST

October 10, 2022

Gu Zhenhao

** Partly adopted from tutorial slides by [Wang Zhi Jian](#).*

Ordered Map ADT

Why do we need the Ordered Map ADT?

- **Purpose:** keep both searching and comparing fast.
- **Operations:**
 1. `insert(i), delete(i)`.
 2. `find(i)`: Check if a key exists.
 3. `predecessor(i), successor(i)`: find the keys next to `i` in the sorted list.
 4. ...

We assume all keys are unique.

Implementation of Ordered Map

- As **unsorted array**:

6	4	9	2	5	8
---	---	---	---	---	---

- Finding a key is very slow.

insert	find	predecessor
$O(1)$	$O(n)$	$O(n)$

Implementation of Ordered Map

- As **sorted array**:

2	4	5	6	8	9
---	---	---	---	---	---

- Can use **binary search** to find keys, faster to find predecessor/successors.
- Slower insertion (need a step in **insertion sort**).

insert	find	predecessor
$O(n)$	$O(\log n)$	$O(1)$

Implementation of Ordered Map

- As **hash map**:

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]
		2		4	5	6		8	9

- Fast insertion/deletion and finding.
- Need to check all keys for predecessor/successor!

insert	find	predecessor
$O(1)$	$O(1)$	$O(n)$

Operations	Unsorted Array	Sorted Array	Hash Map
insert	$O(1)$	$O(n)$	$O(1)$
delete	$O(n)$	$O(n)$	$O(1)$
find	$O(n)$	$O(\log n)$	$O(1)$
predecessor	$O(n)$	$O(1)$	$O(n)$

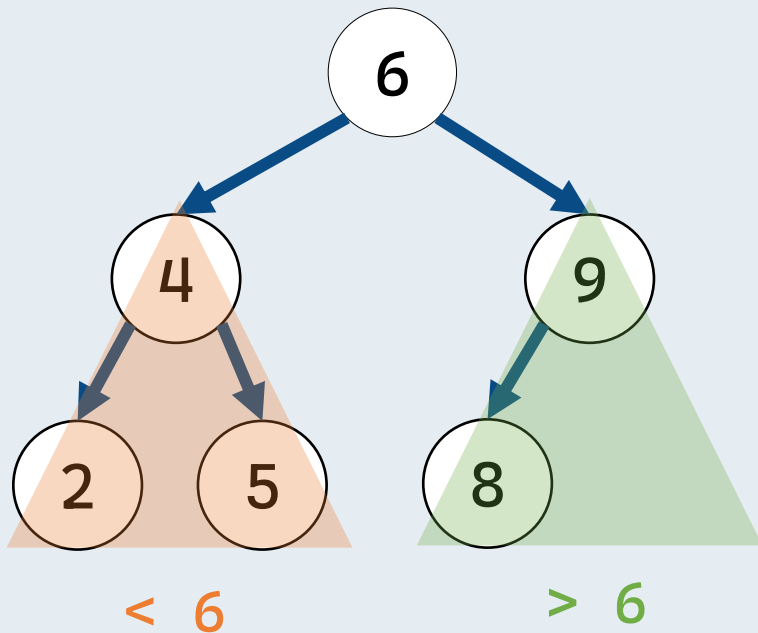
Better at **comparing**

Better at **searching/insert**

- Is there a way to take the merits of these two?

Implementation of Ordered Map

- As binary tree:



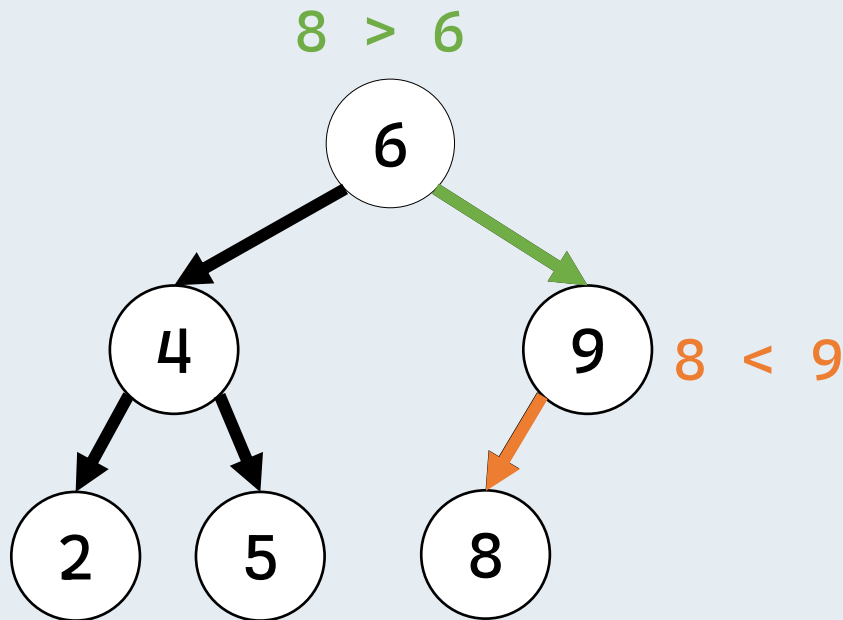
For **every node**:

Each node in **left subtree** is smaller than it,

Each node in the **right subtree** is larger than it.

Implementation of Ordered Map

- `find(8)`:



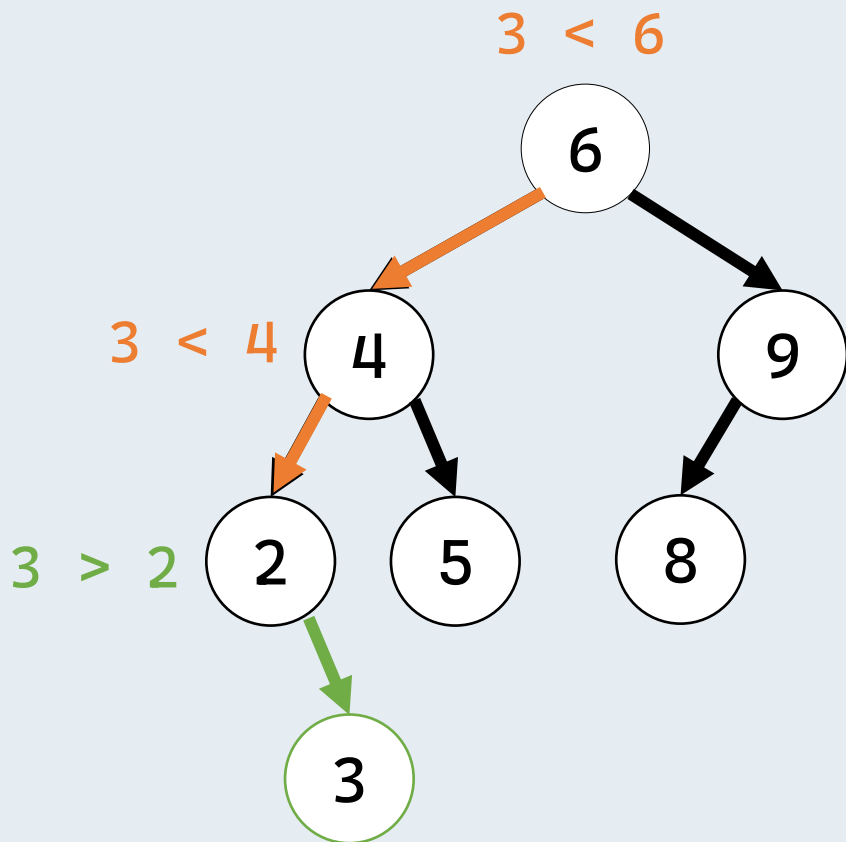
Kind of like
binary search!

To **find** a key:

- Start from **root**,
- If the node is null, the key is not found.
- If the searched key is larger than root, search **right subtree**.
- Otherwise search **left subtree**.

Implementation of Ordered Map

- `insert(3):`



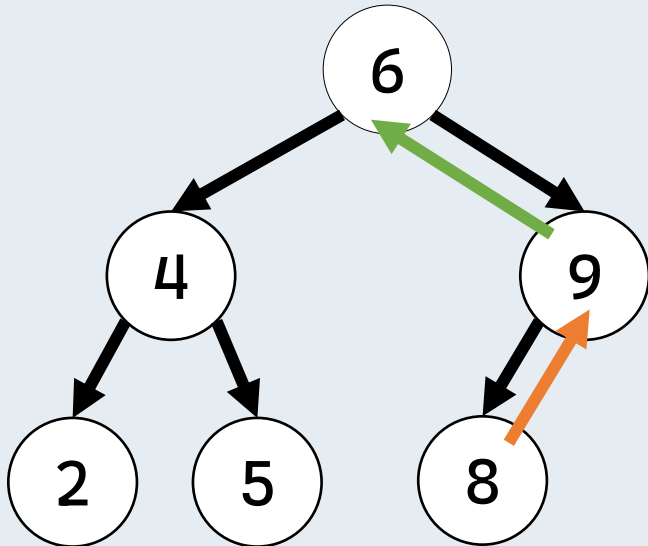
To **insert** a key:

- First find the key.
- If it doesn't exist, insert as a leaf node.

Implementation of Ordered Map

- `predecessor(8)`:

6 is predecessor

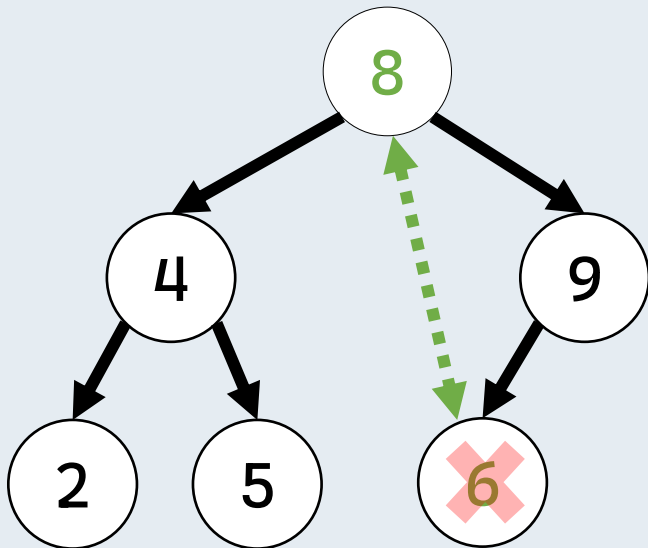


To find **predecessor** of a key:

- If it has a left child, find max in left subtree.
- Otherwise
 - if it is a right child, the predecessor is its parent.
 - Otherwise the predecessor is its “first left ancestor”.

Implementation of Ordered Map

- `delete(6)`:



To **delete** a key:

- If it is a leaf, simply remove the node.
- If it has only one child, remove the node and attach the child to the node's parent.
- If it has 2 children, swap the node with its successor, then remove.

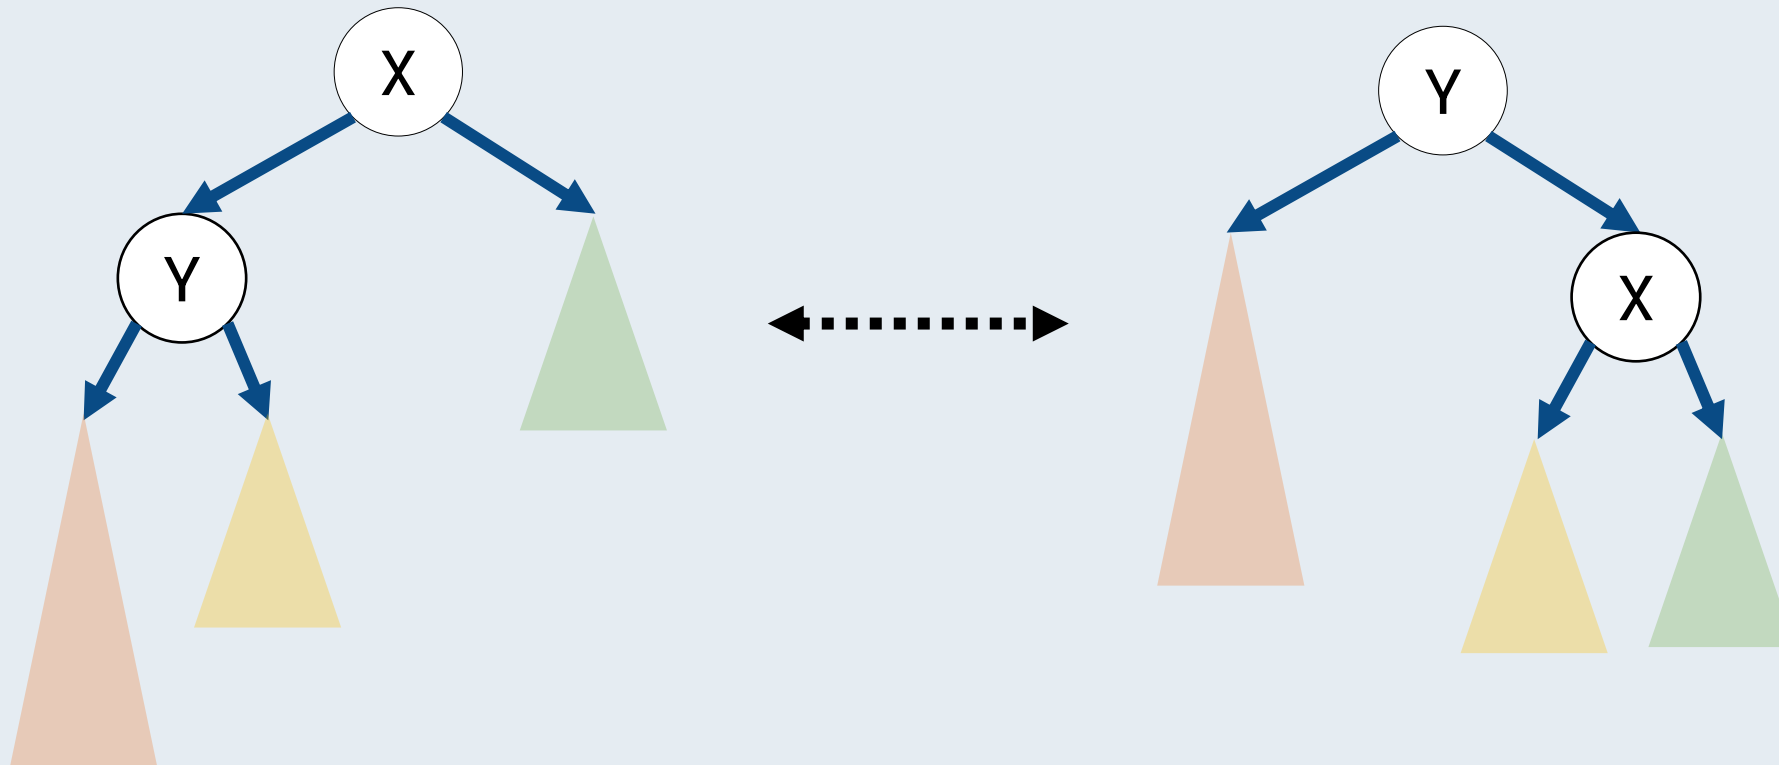
Operations	Sorted Array	Hash Map	BST
insert	$O(n)$	$O(1)$	$O(h)$
delete	$O(n)$	$O(1)$	$O(h)$
find	$O(\log n)$	$O(1)$	$O(h)$
predecessor	$O(1)$	$O(n)$	$O(h)$

How to reduce the height of tree in worst case?

Problem:

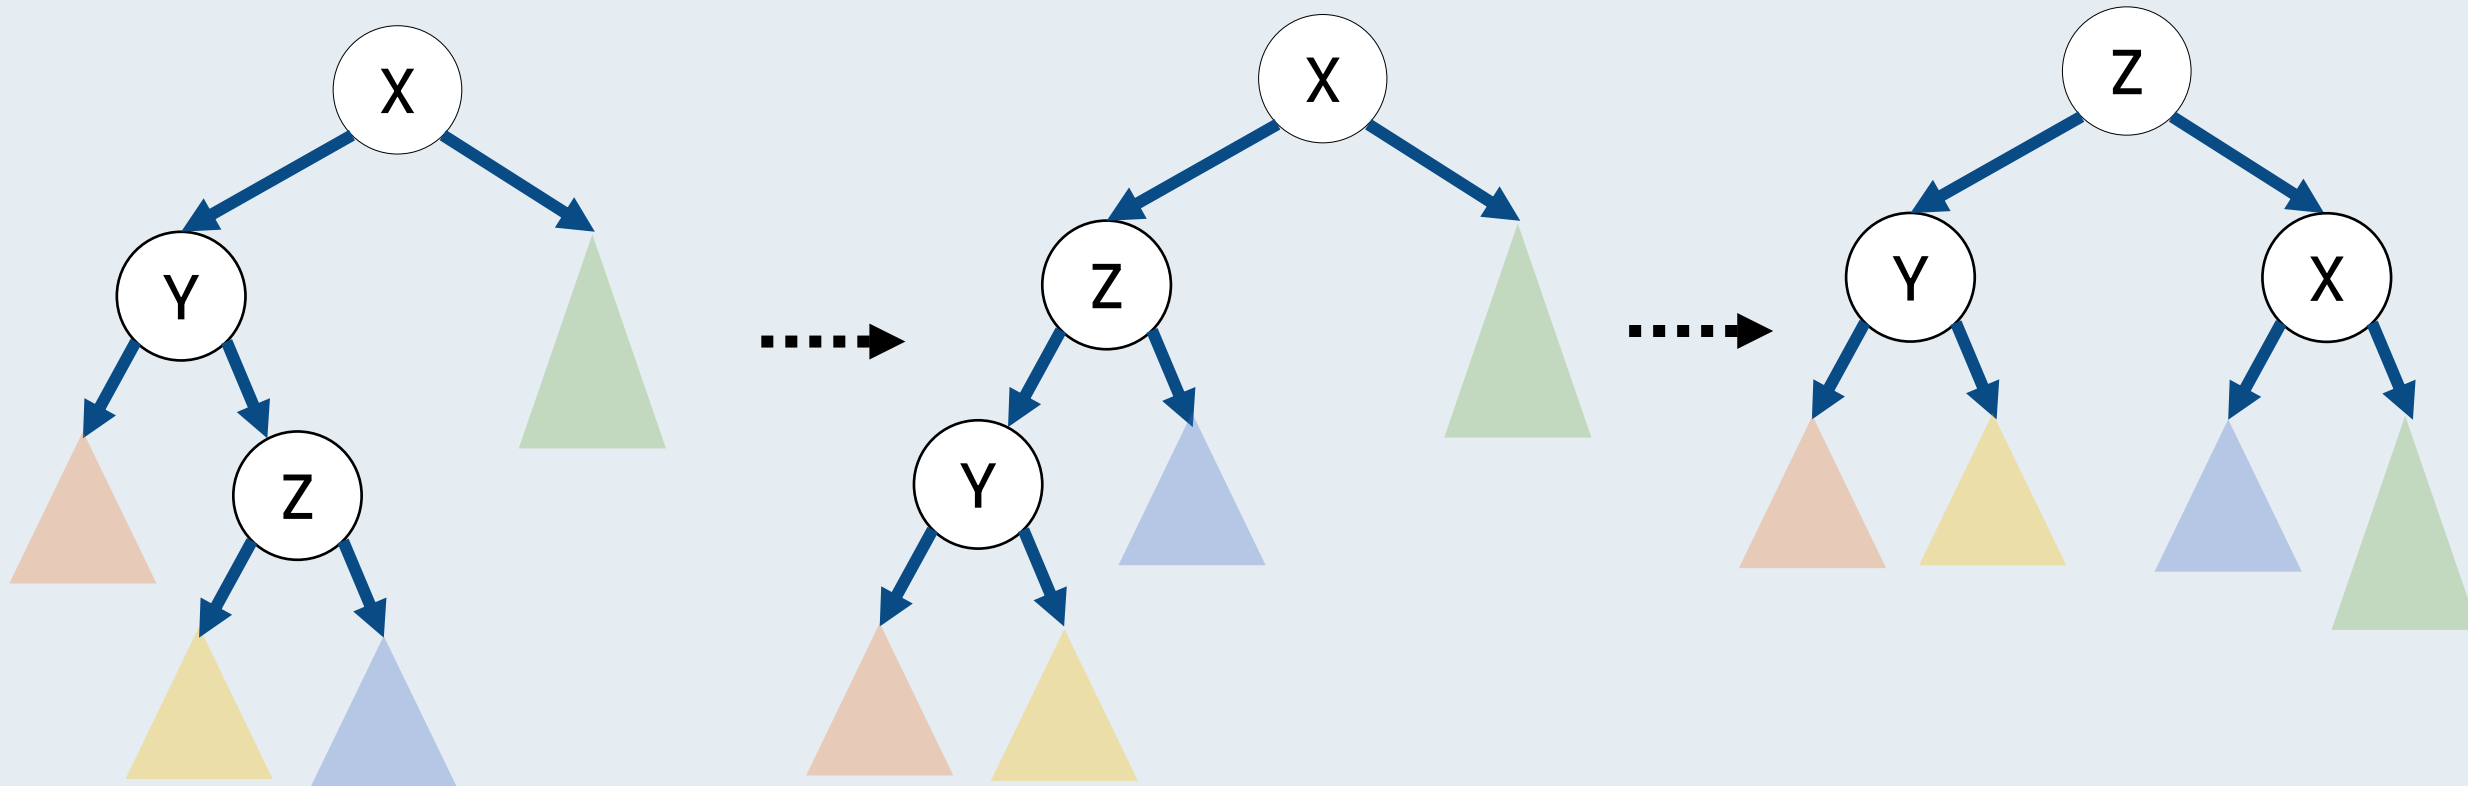
- The height of the tree h can be $O(n)$ in the worst case!

- When tree is unbalanced (has a node where the height of left subtree and right subtree differ more than 1), needs rebalancing!
- **Single rotations:**

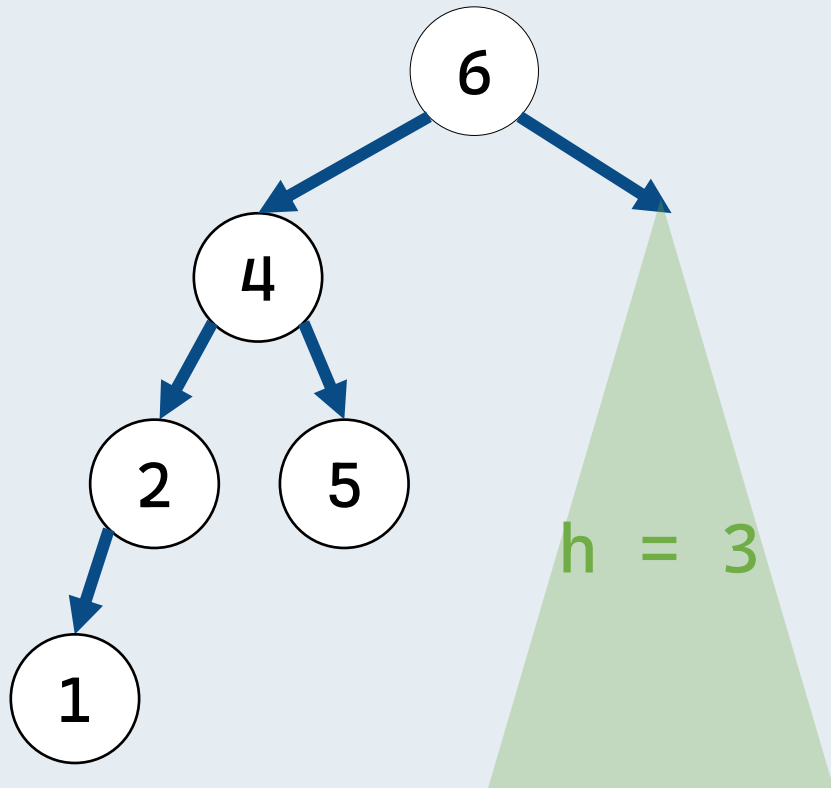


Rebalancing

- When tree is unbalanced (has a node where the height of left subtree and right subtree differ more than 1), needs rebalancing!
- **Double rotations:**



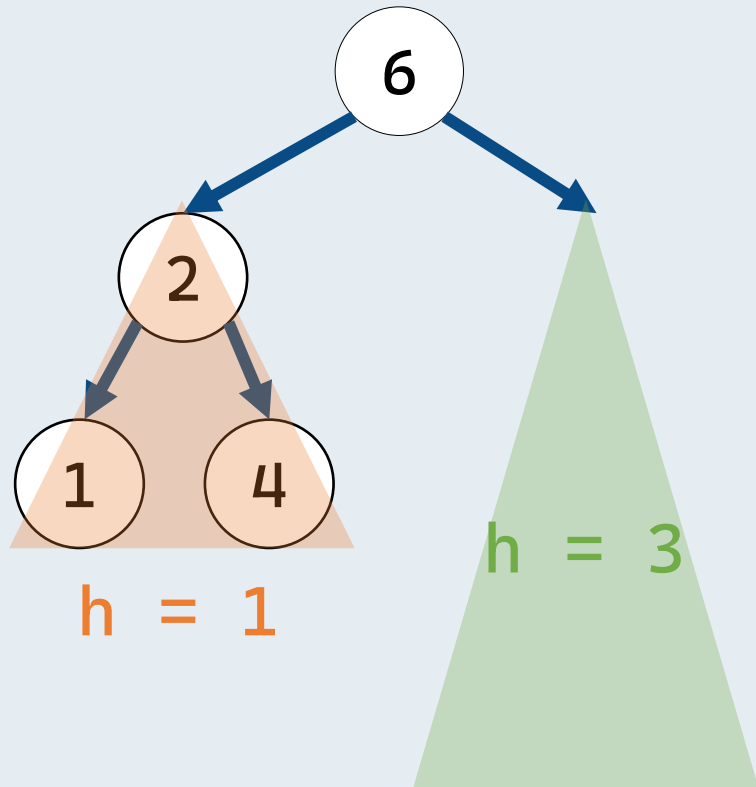
True or false: *Given any AVL tree of height 4, deleting any vertex in the tree will not result in more than 1 rebalancing operation (not rotation but rebalancing operations!).*



False. We can find a counter example:
if we delete 5,

- Need rebalancing in the left subtree,

True or false: *Given any AVL tree of height 4, deleting any vertex in the tree will not result in more than 1 rebalancing operation (not rotation but rebalancing operations!).*



False. We can find a counter example:
if we delete 5,

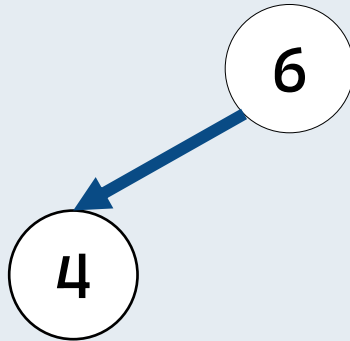
- Need rebalancing in the left subtree,
- Need another to balance the left and right subtree.

True or false: *The minimum number of vertices in an AVL tree of height 5 is 21.*

6

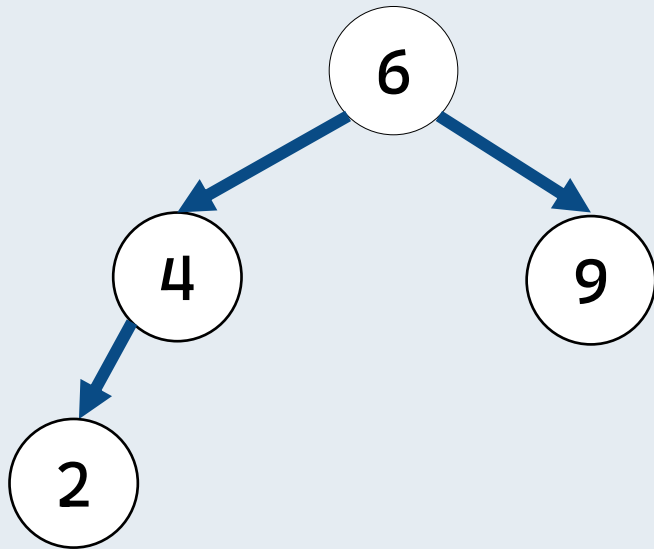
Height of AVL tree	Minimum no. of nodes
0	1
1	
2	
3	
4	
5	

True or false: *The minimum number of vertices in an AVL tree of height 5 is 21.*



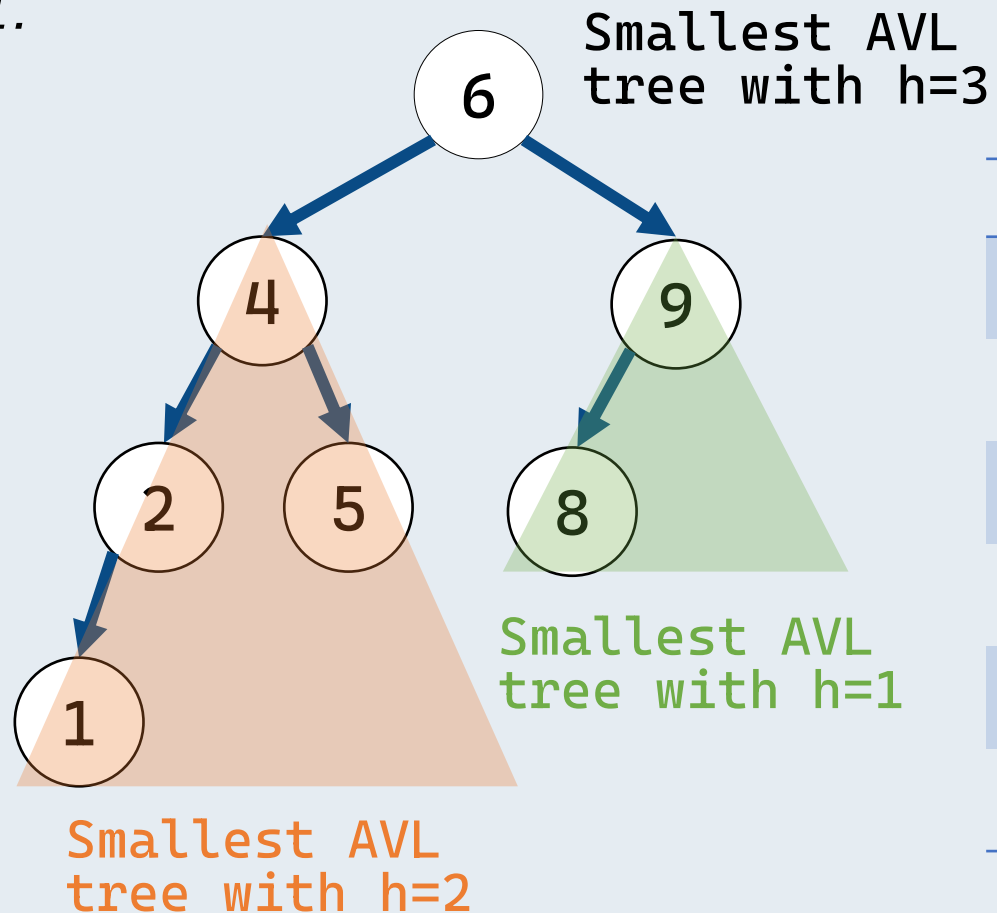
Height of AVL tree	Minimum no. of nodes
0	1
1	2
2	
3	
4	
5	

True or false: *The minimum number of vertices in an AVL tree of height 5 is 21.*



Height of AVL tree	Minimum no. of nodes
0	1
1	2
2	4
3	
4	
5	

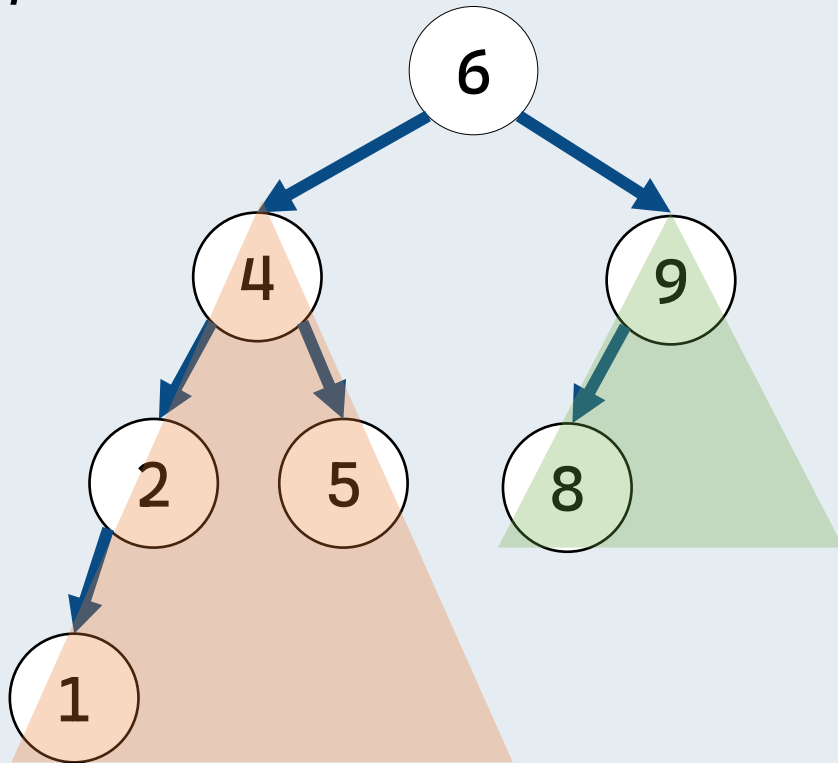
True or false: *The minimum number of vertices in an AVL tree of height 5 is 21.*



Height of AVL tree	Minimum no. of nodes
0	1
1	2
2	4
3	7
4	
5	

$$N_i = N_{i-1} + N_{i-2} + 1$$

True or false: *The minimum number of vertices in an AVL tree of height 5 is 21.*

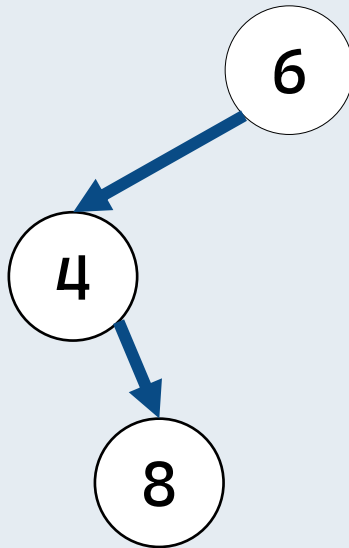


False. The minimum is 20.

Height of AVL tree	Minimum no. of nodes
0	1
1	2
2	4
3	7
4	12
5	20

$$N_i = N_{i-1} + N_{i-2} + 1$$

True or false: *In a tree, if for every vertex x that is not a leaf, $x.\text{left.key} < x.\text{key}$ if x has a left child and $x.\text{key} < x.\text{right.key}$ if x has a right child, the tree is a BST.*



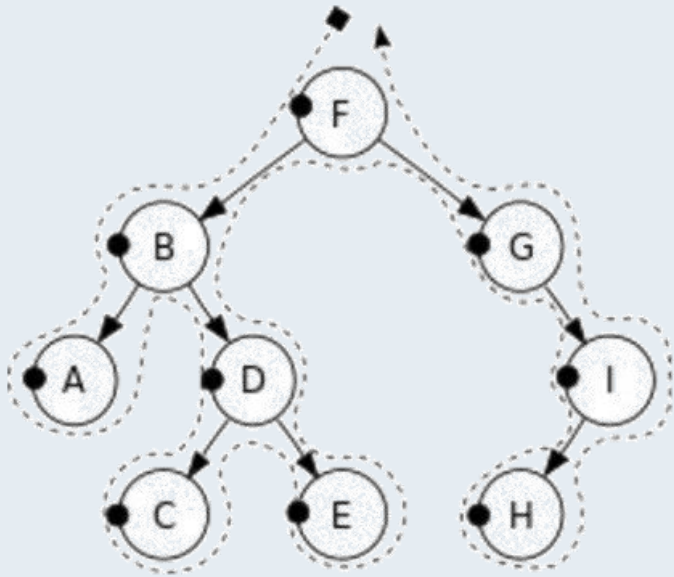
False. We can find a counter example:

The tree on the left is not a BST because we require **ALL** nodes in left subtree are smaller than root.

Tree Traversal

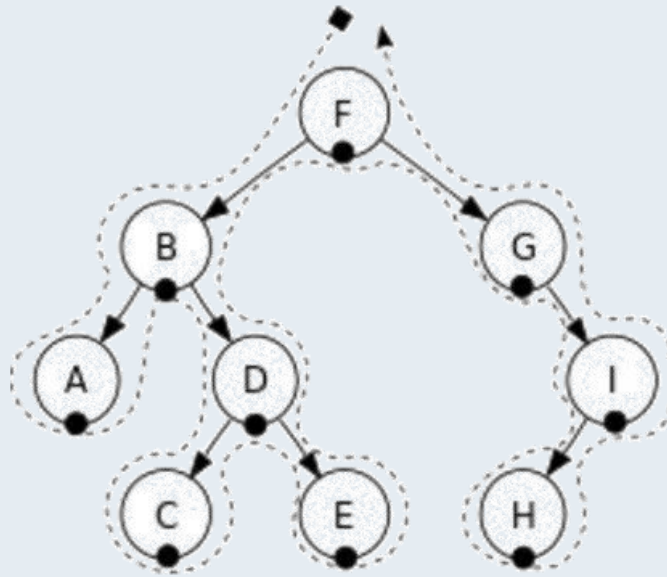
How to properly do traversal?

Three ways of tree traversal



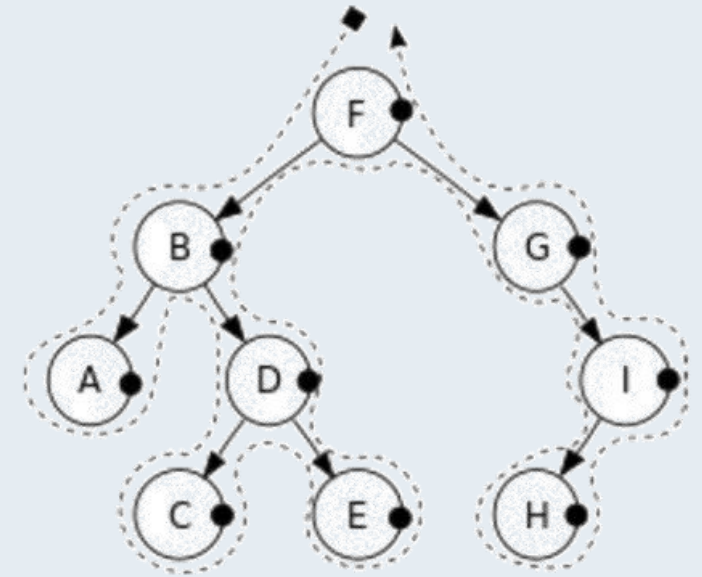
pre-order(node)

```
if node is null: return  
visit(node)  
pre-order(node.left)  
pre-order(node.right)
```



in-order(node)

```
if node is null: return  
in-order(node.left)  
visit(node)  
in-order(node.right)
```



post-order(node)

```
if node is null: return  
post-order(node.left)  
post-order(node.right)  
visit(node)
```

Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$                     Start from smallest node in BST...
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$ 
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$ 
6:   end while                                                  ▷ successor returns -1 if there is no successor
7: end procedure
```

What is the running time of Algorithm 1?

Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$ 
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$     Visit current node...
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$ 
6:   end while                                                    ▷ successor returns -1 if there is no successor
7: end procedure
```

What is the running time of Algorithm 1?

Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$ 
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$                                 Go to next smallest node,
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$  until all nodes are visited.
6:   end while                                                ▷ successor returns -1 if there is no successor
7: end procedure
```

What is the running time of Algorithm 1?

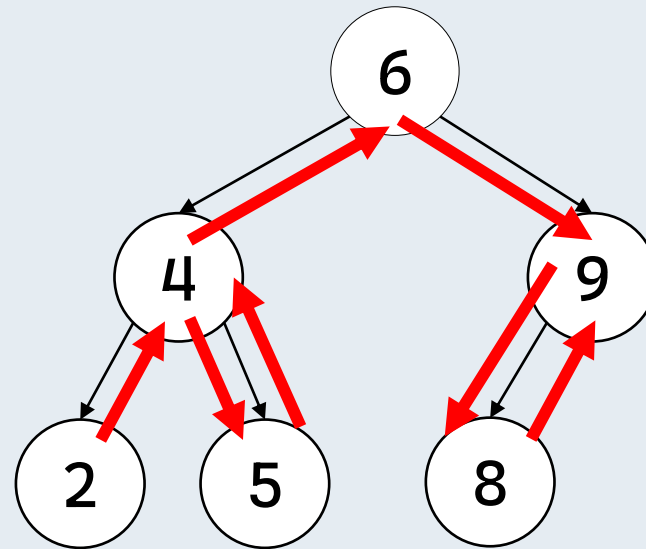
Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$                      $O(\log n)$  time.
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$                                  $O(\log n)$  time, the while
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$  loop runs  $O(n)$  times.
6:   end while                                                ▷ successor returns -1 if there is no successor
7: end procedure
```

What is the running time of Algorithm 1?

In total $O(n \log n)$ time.

Propose modifications to the **successor** function such that Algorithm 1 runs in $O(n)$ time.



The effect is exactly the same as in-order traversal...

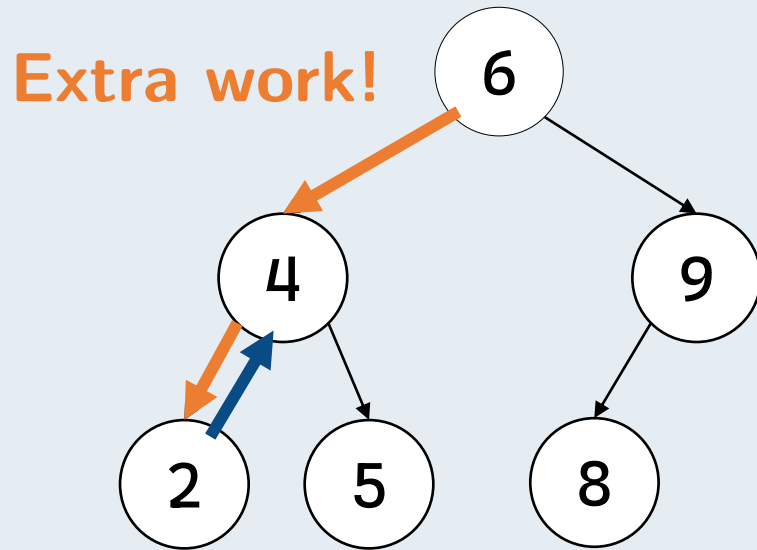
But in-order traversal runs in $O(n)$ time. **What's the difference?**

Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$ 
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$ 
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$ 
6:   end while                                                    ▷ successor returns -1 if there is no successor
7: end procedure
```

We need to first find the node containing **currentNodeValue** from the root node, then go to its successor!

If we call `successor(2)`...



1. Find the node containing 2...
2. Starting from the node, find its successor.

Algorithm 1 In-Order Traversal

```
1: procedure INORDERTRAVERSAL( $T$ )                                ▷ bBST  $T$  is supplied as input
2:    $currentNodeValue \leftarrow T.findMin()$ 
3:   while  $currentNodeValue \neq -1$  do
4:     output  $currentNodeValue$ 
5:      $currentNodeValue \leftarrow T.successor(currentNodeValue)$ 
6:   end while                                                    ▷ successor returns -1 if there is no successor
7: end procedure
```

Idea: We feed a reference of **currentNode** to the successor function instead, instead of its value.

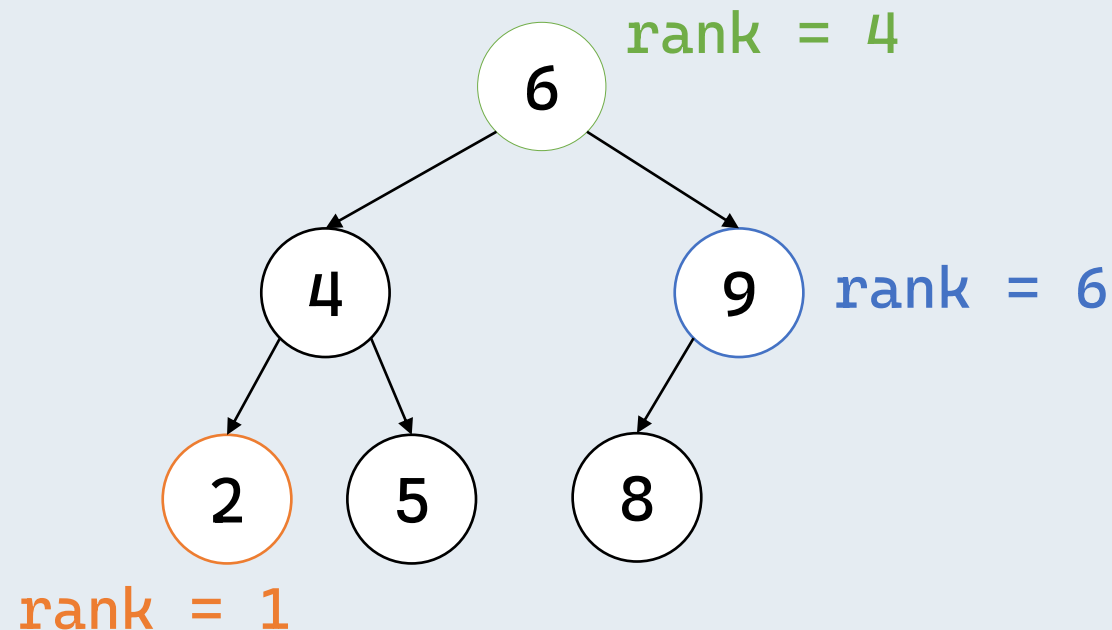
Each node is now visited at most $O(1)$ times ... $O(n)$ time in total!

Applications

How to use BST's ability of comparing keys?

A node x has **rank** k in a BST if there are $k - 1$ nodes that are smaller than x in the BST.

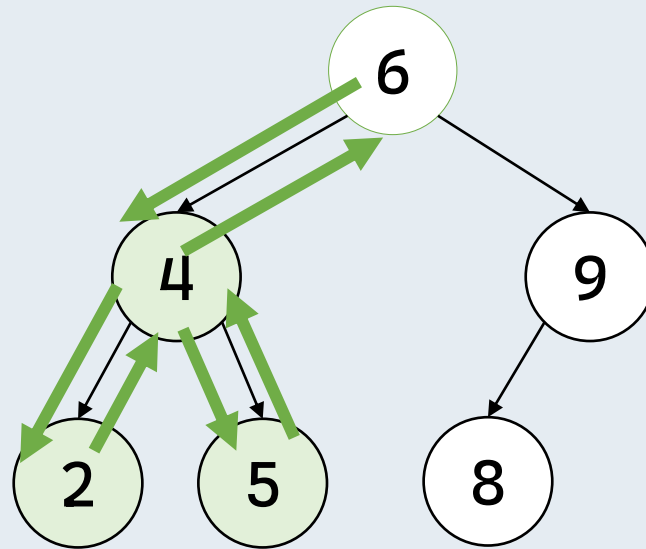
Goal: find the rank of a node x in $O(h)$ time.



Trivial Solution: Find all nodes that are smaller than the node x .

We can use in-order traversal, count the visited nodes, until we visit x !

Time needed: $O(n)$ time.

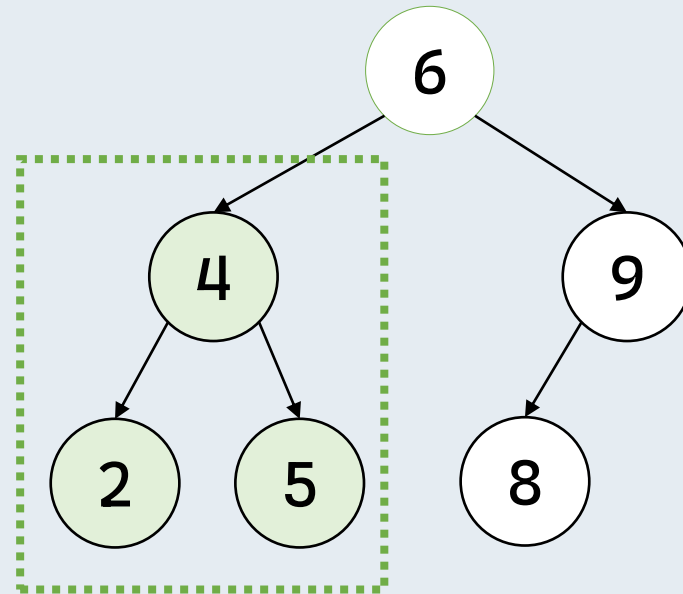


Where are the nodes smaller than the root 6 located?

The nodes that are smaller than root are all in left subtree.

Idea: just return the size of the left subtree!

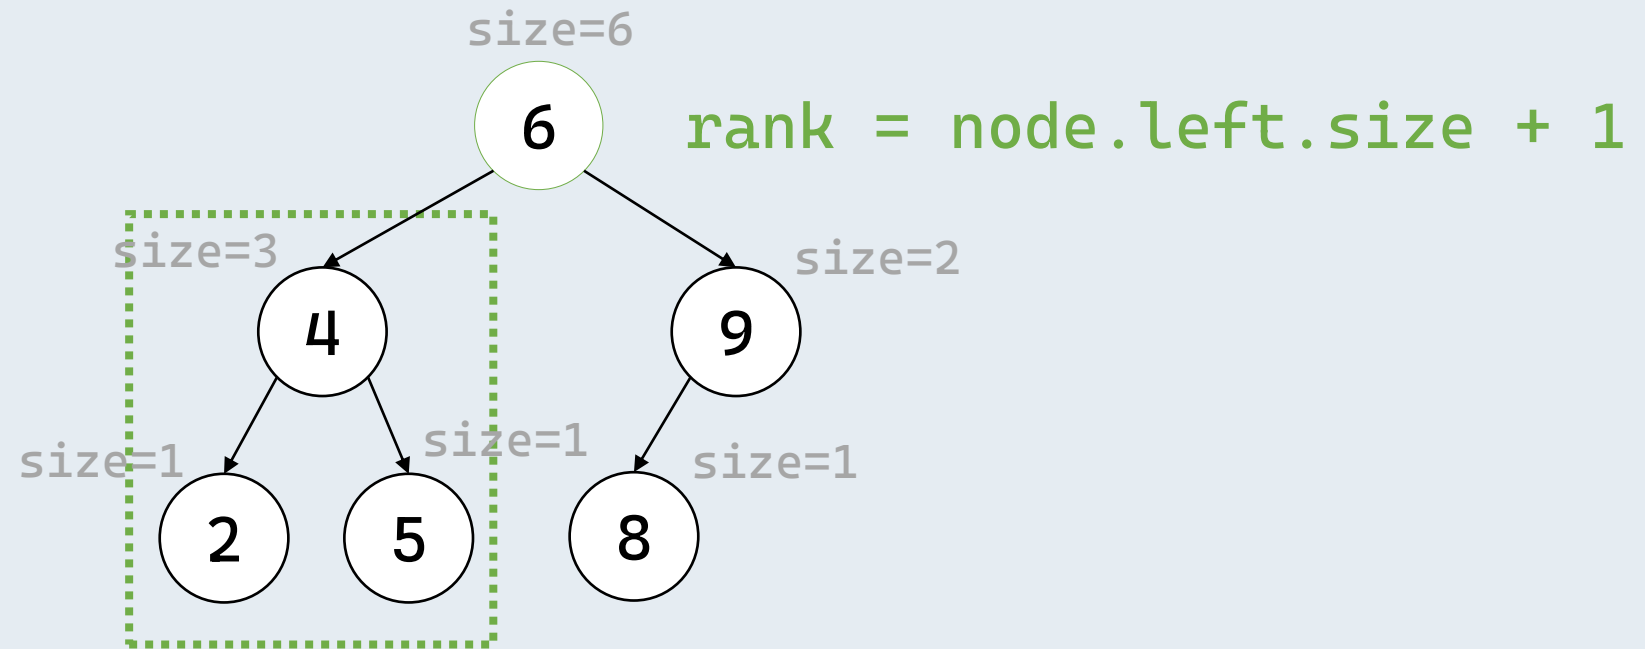
But counting the size is still $O(n)$...



The nodes that are smaller than 6 are all in left subtree.

Idea: just return the size of the left subtree!

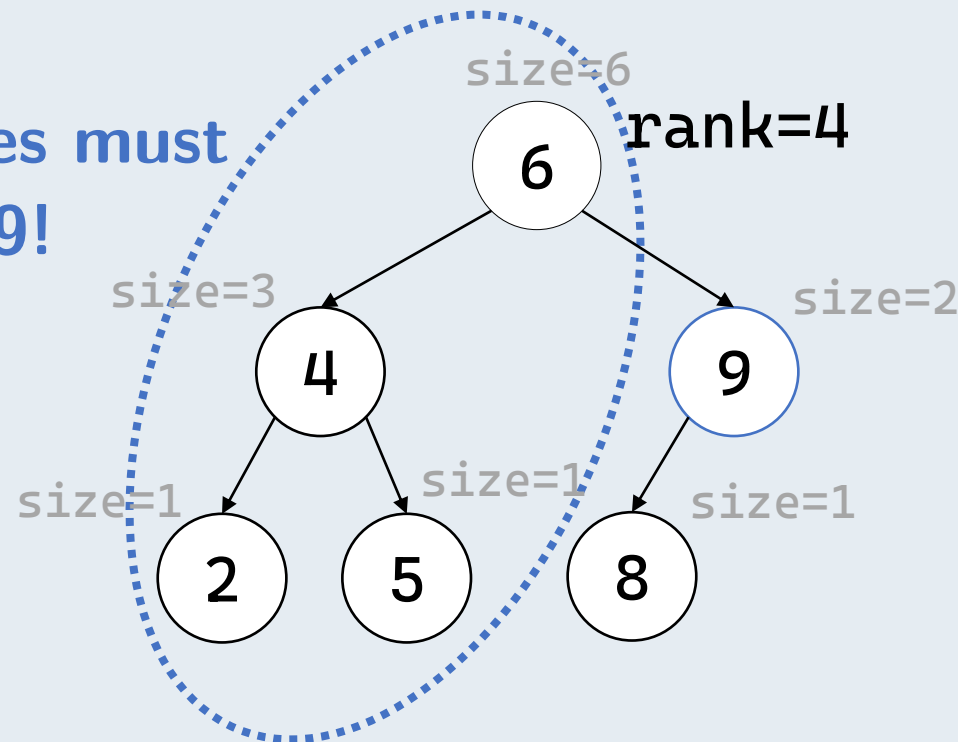
Count and store the **size** of subtree rooted at the node beforehand!



What if we want to find the rank of nodes other than root, e.g. 9?

1. Visit the root, and the root $6 < 9$.

All of these nodes must be smaller than 9!



- We know that 6 is in rank 4 in sorted array ... And all numbers before 6 are smaller than 6.

4	2	5	6	9	8
---	---	---	---	---	---

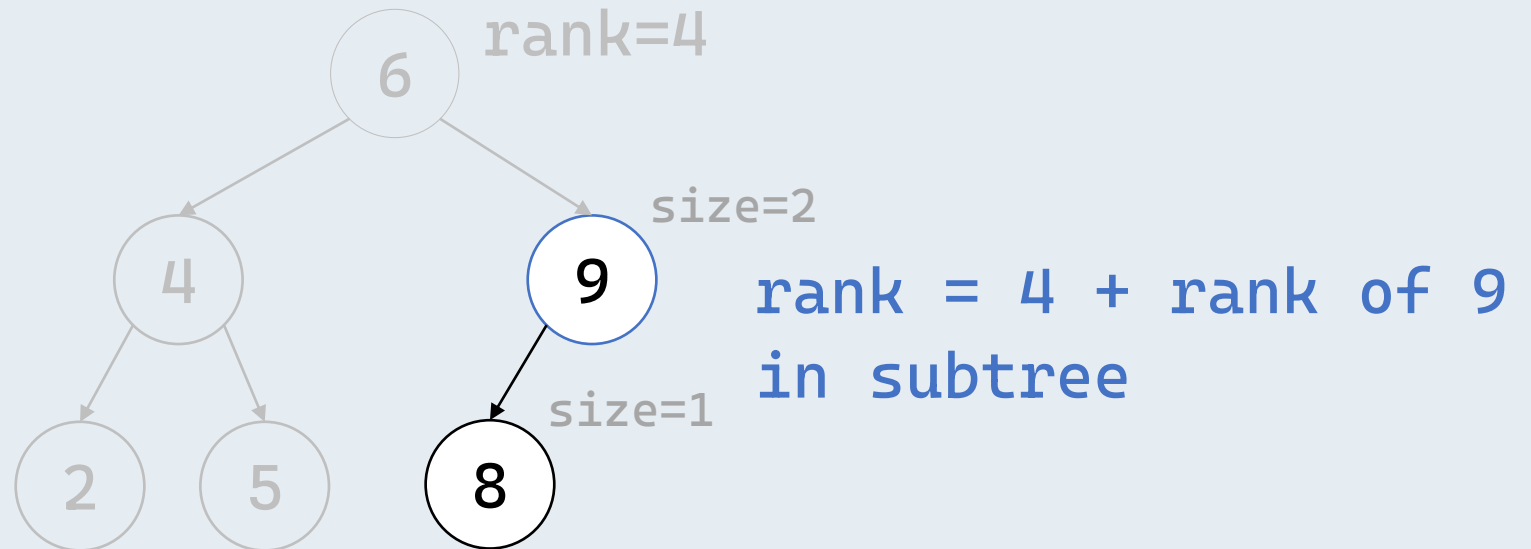
rank=4

- rank of 9 in the whole array = rank of 6 + rank of 9 in right sub-array.

4	2	5	6	9	8
---	---	---	---	---	---

What if we want to find the rank of node other than root, e.g. 9?

1. Visit the root, and the root $6 < 9$.
2. Ignore the left subtree, and find the rank of 9 in right subtree.



Algorithm 2 BST Rank

```
1: function RANK(node, v)
2:   if node.key = v then
3:     return node.left.size + 1
4:   else if node.key > v then
5:     return RANK(node.left, v)
6:   else
7:     return node.left.size + 1 + RANK(node.right, v)
8:   end if
9: end function
```

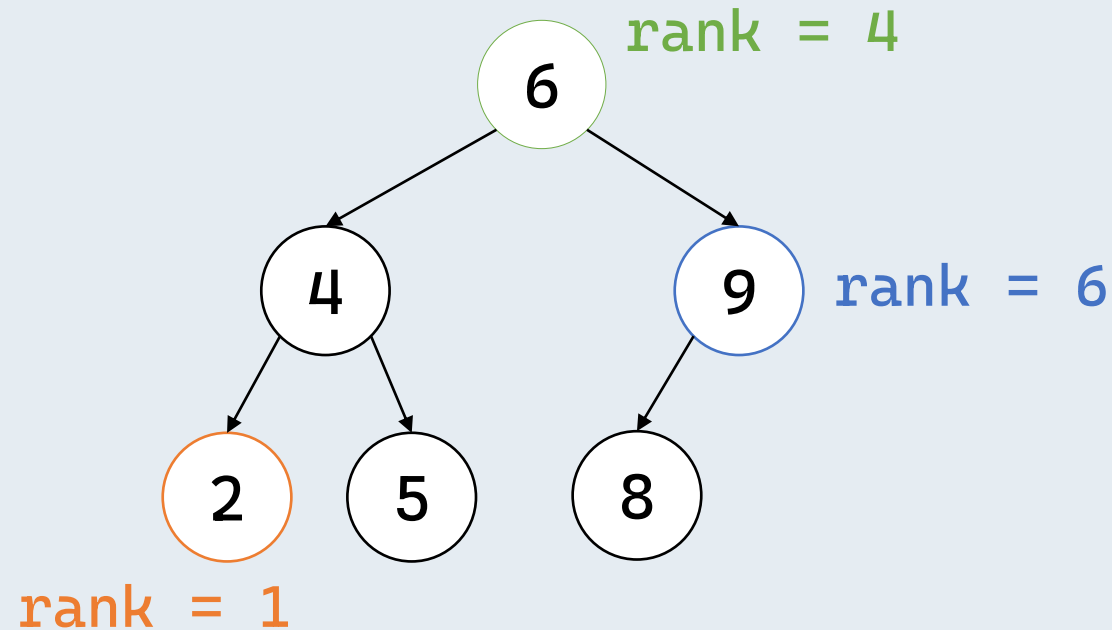
Recurrence relation:

$$T(h) = T(h - 1) + 1$$

In total $O(h)$ time.

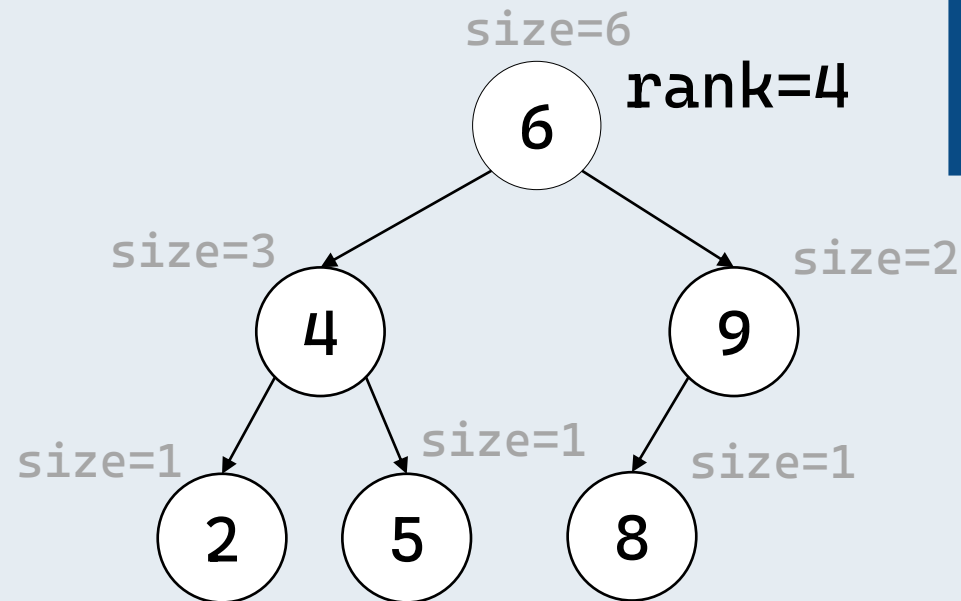
A node x has **rank** k in a BST if there are $k - 1$ nodes that are smaller than x in the BST.

Goal: Find the node with rank k in $O(h)$ time.



Example: Find the node with rank = 5.

$5 > \text{rank of root node} = 4.$

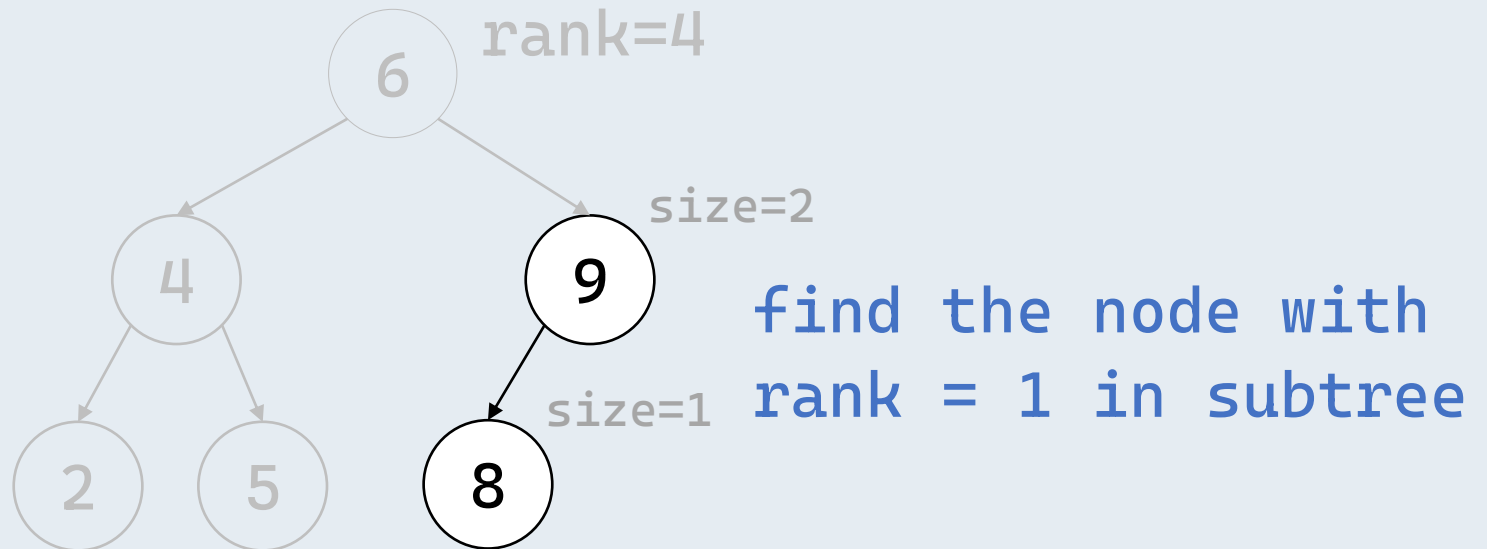


Do we need to check every node for their rank?

Example: Find the node with rank = 5.

5 > rank of root node = 4.

Idea: we only need to check the right subtree for the node!



Algorithm 3 BST Select

```
1: function SELECT(node, k)
2:    $q \leftarrow \text{node.left.size}$ 
3:   if  $q + 1 = k$  then
4:     return node.key
5:   else if  $q + 1 > k$  then
6:     return SELECT(node.left, k)
7:   else
8:     return SELECT(node.right,  $k - q - 1$ )
9:   end if
10: end function
```

This is very similar to
QuickSelect!

... Except for the
partition is already
done by BST!

** It is very common to use recursions in tree-related algorithms! Get used to it ☺*

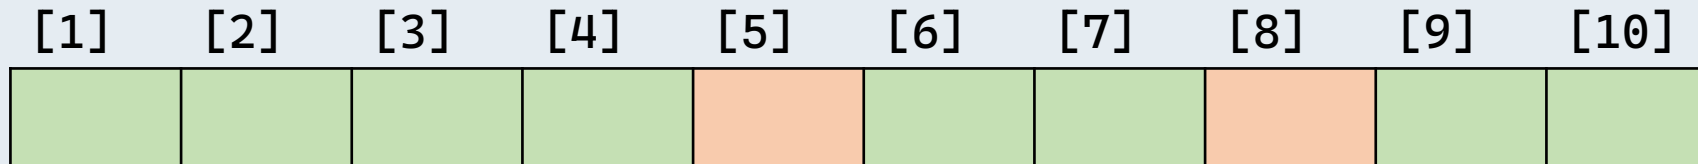
- n servers in a line, either enabled or disabled.
- Server i can send data to server j if all servers between i and j are enabled.

Goal: support the operations:

1. `enable(i)`: Enable the i -th server.
2. `disable(i)`: disable the i -th server.
3. `send(i, j)`: return true if we can send data from server i to j .

- n servers in a line, either enabled or disabled.
- Server i can send data to server j if all servers between i and j are enabled.

Example: Servers 5 and 8 are disabled.



`send(3, 10)` returns `false`.

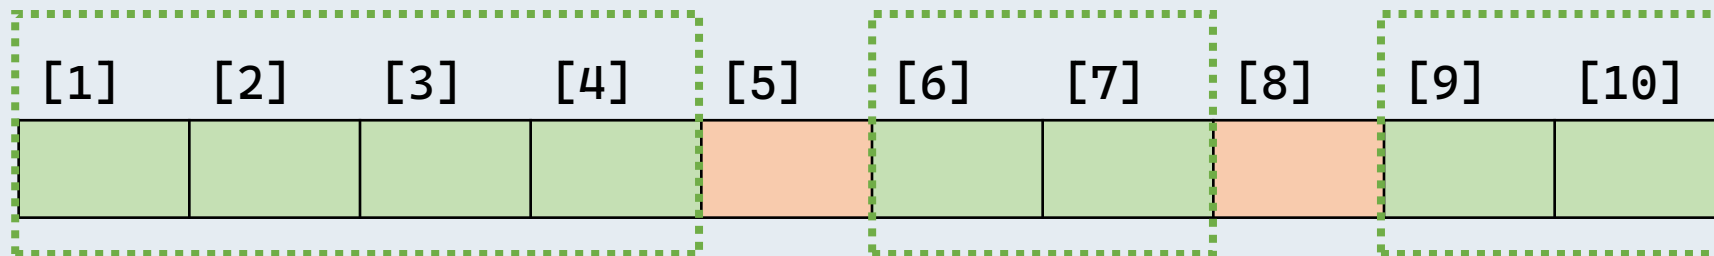
Trivial Solution:

- Store everything in an array.
- Start from **i** and see if we encounter a disabled server, until we reach **j**.
- **enable** and **disable** are $O(1)$ time, **send** takes $O(n)$ time.

[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]	[10]
T	T	T	T	F	T	T	F	T	T

Another Solution:

- Store everything in UFDS.
- Group consecutive enabled servers in a set.
- **send** simply checks if **i** and **j** are in same set.
- **enable** and **send** are $O(\alpha(n))$ time, **disable** takes $O(n)$ time.

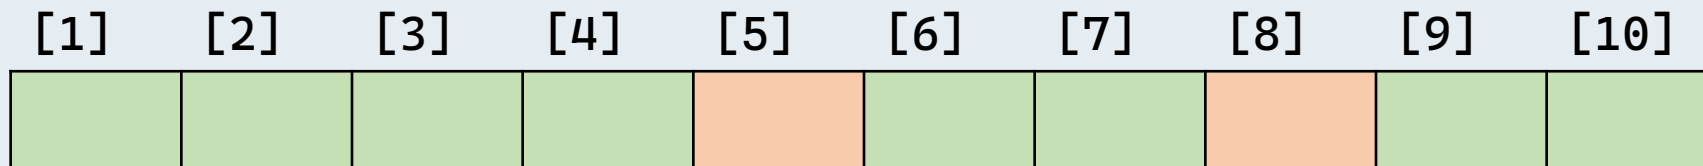


Observation:

Cannot send message from i to j if the first disabled server after i is $< j$.

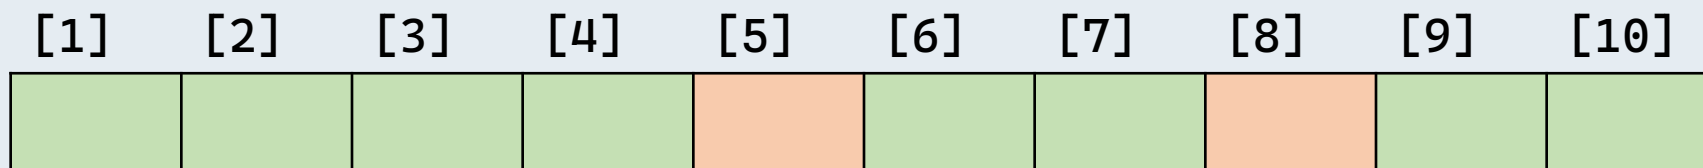
e.g. `send(3, 10)` is false because the first disabled server after 3 is 5.

find the **successor** of 3!



Idea:

- Store all the **disabled servers** in an AVL tree.
- **send(*i*, *j*)** checks existence of ***i***, and find the successor of ***i*** and check if it is $> j$.
- All 3 operations takes $O(\log n)$ time.



Algorithm 4 Solution to Problem 4

```
1:  $T \leftarrow$  bBST, initially empty
2: procedure ENABLE( $i$ )
3:    $T.delete(i)$ 
4: end procedure
5: procedure DISABLE( $i$ )
6:    $T.insert(i)$ 
7: end procedure
8: function SEND( $i, j$ )
9:   if  $i$  is in  $T$  or  $T.successor(i) \leq j$  then
10:    return false
11:  else
12:    return true
13:  end if
14: end function
```

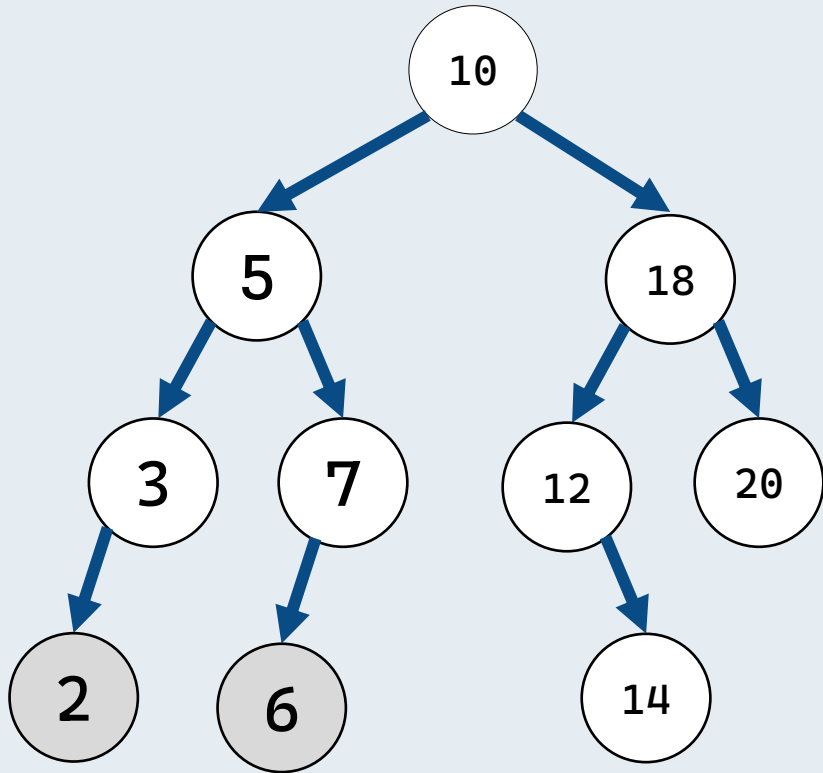
Note:

This **successor** is different from the one in the lecture: **i** may not be a node in the BST!

How to implement this successor function?

** See the appendix for hints!*

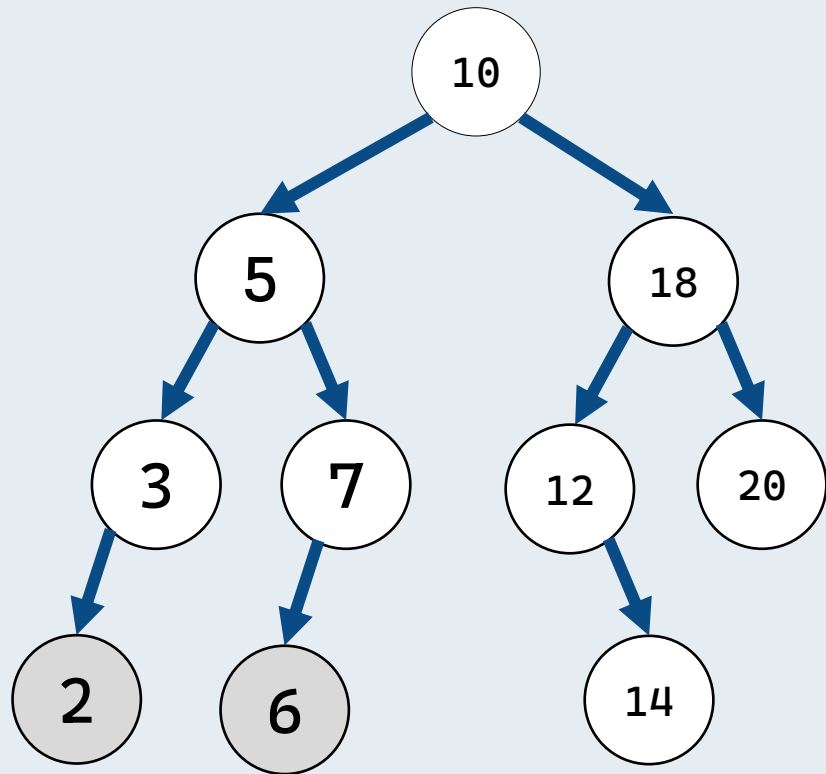
The Lowest Common Ancestor (LCA) of two nodes a and b in a BST is the node furthest from the root that is an ancestor of both a and b .



Example:

LCA of 2 and 6 is 5.

The Lowest Common Ancestor (LCA) of two nodes a and b in a BST is the node furthest from the root that is an ancestor of both a and b .



Idea:

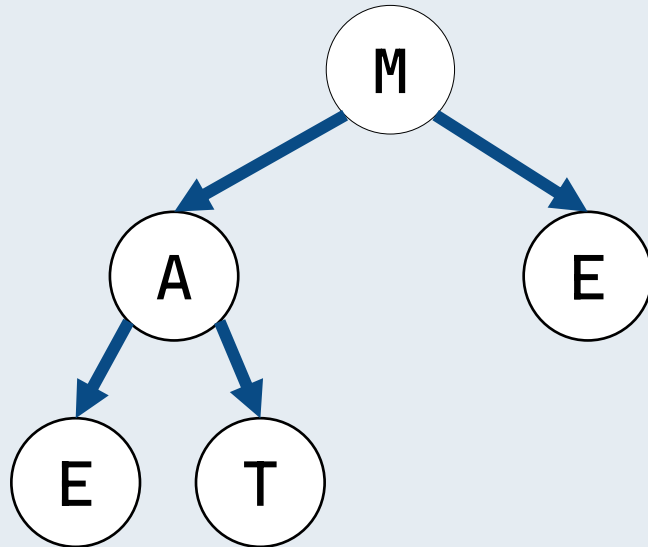
Do we need to store the path?

- We can first find the two nodes:
 - For node 2, we go through $10 \rightarrow 5 \rightarrow 3 \rightarrow 2$.
 - For node 6, we go through $10 \rightarrow 5 \rightarrow 7 \rightarrow 6$.
- **Just check the common path!**

A set of N strings of length L .

Goal: implement `gotPrefix(k)`: return true if string k is a prefix of some string in the set.

Example: Set is $\{ME, MAE, MAT\}$, `gotPrefix("MA")` returns **true**.



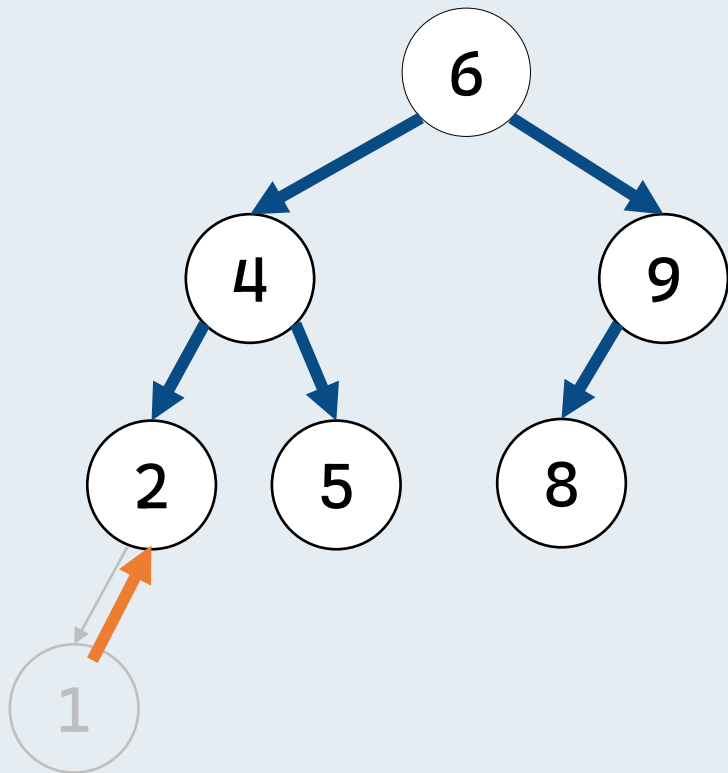
Idea: Simply use a trie!

- **preprocessing:** each insertion takes $O(L)$ time.
- **query:** check if there is a path matching k . Takes $O(k.length)$ time.

Appendix

successor(value)

Problem: How do we implement `successor(value)`, where the value might not be among the nodes in BST?



Example: `successor(1)`.

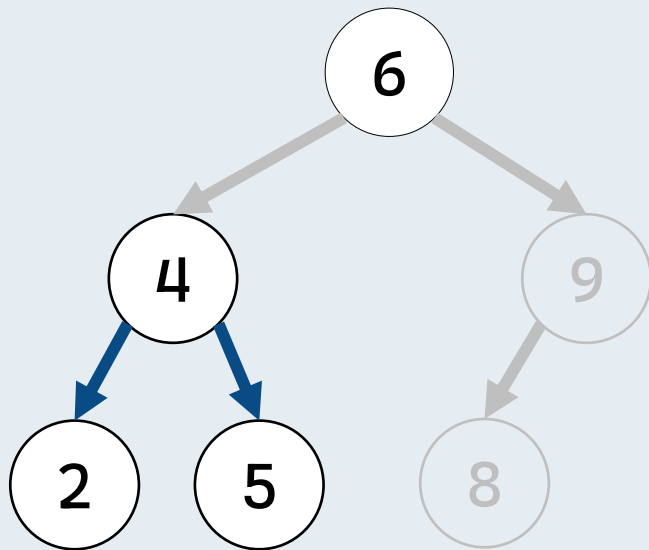
Trivial Solution:

- Insert a node 1,
- Find its successor,
- Delete the node 1.

Can we avoid creating the node?

successor(value)

Problem: How do we implement $\text{successor}(\text{value})$, where the value might not be among the nodes in BST?



Idea: Can also do it recursively!

We have the root $6 > 1$.

If 6 is not the successor, it must be in left subtree!

The successor must be: $\min\{6, \text{successor of 1 in left subtree}\}$.

End of File

Thank you very much for your attention :-)